



Predicting Real or Fake Job Posting Using Machine Learning

By Buvana
Mini Project 3

01

Data Cleaning &
Investigation

03

Text Data Cleaning &
Count Vectorization

05

Results and Conclusion

02

EDA & Feature
Extraction

04

Machine Learning Models-text
and feature selected variables

Problem Statement

Predict the **probabilities** of fake job postings.

Objective

To build machine learning models using data to predict which **job postings are real or fake** and to draw insightful data.


Scope

Use of models to perform **text data classification and to deploy models on the features selected from forward feature selection.**
Use of accuracy metrics to evaluate the model results

- ❖ Fake job advertisements found in recruitment portals –can go detected or undetected.
- ❖ Exponential increase in fraudulent job opportunities
- ❖ In USA- Spike in fraud job adverts/ spams in Q2 of 2020 (Source: CNBC)
- ❖ **For jobseekers: Waste of time and effort applying for non existent jobs**
- ❖ **For companies/recruiters : Reputation and credibility**

Fraud reports about business and job-related opportunities

By subcategory, quarterly since 2015



SOURCE: Federal Trade Commission

benefits	telecommuting	has_company_logo	has_questions	employment_type	required_experience	required_education	industry	function	fraudulent
NaN	0	1	0	Other	Internship	NaN	NaN	Marketing	0
What you will get from usThrough being part of...	0	1	0	Full-time	Not Applicable	NaN	Marketing and Advertising	Customer Service	0
NaN	0	1	0	NaN	NaN	NaN	NaN	NaN	0
Our culture is anything but corporate—we have ...	0	1	0	Full-time	Mid-Senior level	Bachelor's Degree	Computer Software	Sales	0
Full Benefits Offered	0	1	1	Full-time	Mid-Senior level	Bachelor's Degree	Hospital & Health Care	Health Care Provider	0
NaN	0	0	0	NaN	NaN	NaN	NaN	NaN	0
Your Benefits: Being part of a fast-growing co...	0	1	1	Full-time	Mid-Senior level	Master's Degree	Online Media	Management	0
Competitive Pay. You'll be able to eat steak e...	0	1	1	NaN	NaN	NaN	NaN	NaN	0
NaN	0	1	1	Full-time	Associate	NaN	Information Technology and Services	NaN	0
NaN	0	1	0	Part-time	Entry level	High School or equivalent	Financial Services	Customer Service	0

❖ 17880 Rows


❖ 17 Columns

❖ Text Variables

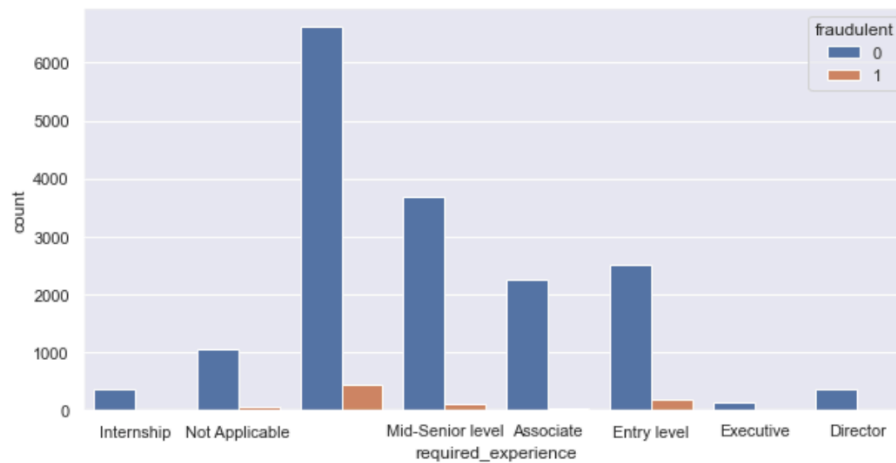
1. Title
2. Company profile
3. Description
4. Requirements
5. Benefits


❖ Target Variable
Fraudulent

EDA



Correlation Matrix





- ❖ Unbalanced data points- Only 4.8% of job postings are fake
- ❖ **Resampling method** to balance out the minority class.
- ❖ Initially 4.8% of the fraudulent columns denoted fake postings. With resampling, we inflated the count to reflect 50% -50% of real and fake job listings

- ❖ Basic feature extraction to the 5 Text Columns to check for any distinct features that could differentiate a real and a fake job posting
- ❖ **Word Count**
- ❖ **Character Count**
- ❖ **Word Density**

char_count_d	char_count_r	char_count_b	char_count_t	char_count_cp	word_density_d	word_density_cp	word_density_r	word_density_t	word_density_b
752	715	0	15	710	6.064516	5.035461	6.217391	7.500000	0.000000
1648	1172	1022	35	909	5.333333	6.060000	6.267380	5.833333	4.542222
299	1173	0	34	728	5.980000	5.352941	7.152439	8.500000	0.000000
2213	1228	669	28	509	6.414493	5.988235	7.057471	5.600000	6.968750
1300	653	19	17	1327	7.142857	6.473171	7.337079	5.666667	6.333333

Cleaning of text variables

- ❖ Removal of Stop words
- ❖ Lemmatized words
- ❖ Tokenization

Count Vectorisation

- ❖ Each of the 5 text columns were fit with Count Vectorizer one by one
- ❖ Transformed each text column and concatenated the data frames of 5 text columns to perform the modelling

Model analysis- Text columns as the predictor columns

- ❖ **Support Vector**
- ❖ **Logistic Regression**
- ❖ **Naïve Bayes Classifier**
- ❖ **Boosting (Gradient Descent)**

	Desc	Desc+Title	Desc+Title+Req	Desc+Title+Req+cp	Desc+Title+Req+cp+Benefits
SVM	0.504195	0.494687	0.500839	0.494966	0.508110
Logreg	0.493568	0.494407	0.501398	0.496085	0.493009
Naive Bayes	0.494128	0.501957	0.497763	0.515660	0.510347
Gradient Boost	0.504474	0.496085	0.493009	0.501957	0.501398

- ❖ The highest accuracy level is only 51.6% when we run Naïve Bayes on all the 4 text columns of description, job title, requirements, company profile.

Forward feature selection resulted:

- ❖ **Predictor columns (X)** of : *'telecommuting', 'has_company_logo', 'has_questions', 'required_experience', 'required_education', 'industry', 'function', 'word_count_t', 'word_count_r', 'char_count_b', 'word_density_cp', 'word_density_r', 'description1', 'title1', 'requirements1', 'benefits1', 'company_profile1'*
(from 34 to 18 variables)
- ❖ **Target Variable (Y):** *'fraudulent'*

Machine Learning Models were rerun

- ❖ **Support Vector**
- ❖ **Logistic Regression**
- ❖ **Naïve Bayes Classifier**
- ❖ **Boosting (Gradient Descent)**

With...

- ❖ **Train-80% Split**
- ❖ **Test-20% Split**
- ❖ **Cross Validation- 10 folds**

Training Set

* SVM *



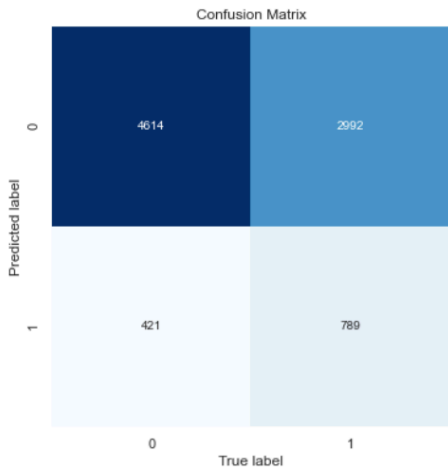
Accuracy : 0.6129 [TP / N] Proportion of predicted labels that match the true labels. Best: 1, Worst: 0

Precision: 0.6521 [TP / (TP + FP)] Not to label a negative sample as positive. Best: 1, Worst: 0

Recall : 0.2087 [TP / (TP + FN)] Find all the positive samples. Best: 1, Worst: 0

ROC AUC : 0.5625 Best: 1, Worst: < 0.5




TP: True Positives, FP: False Positives, TN: True Negatives, FN: False Negatives, N: Number of samples



Test Set

Accuracy : 0.6213 [TP / N] Proportion of predicted labels that match the true labels. Best: 1, Worst: 0
Precision: 0.6512 [TP / (TP + FP)] Not to label a negative sample as positive. Best: 1, Worst: 0
Recall : 0.2202 [TP / (TP + FN)] Find all the positive samples. Best: 1, Worst: 0
ROC AUC : 0.5671 Best: 1, Worst: < 0.5

TP: True Positives, FP: False Positives, TN: True Negatives, FN: False Negatives, N: Number of samples






Training Set

* Logistic *

Accuracy : 0.7459 [TP / N] Proportion of predicted labels that match the true labels. Best: 1, Worst: 0
 Precision: 0.7216 [TP / (TP + FP)] Not to label a negative sample as positive. Best: 1, Worst: 0
 Recall : 0.6636 [TP / (TP + FN)] Find all the positive samples. Best: 1, Worst: 0
 ROC AUC : 0.7357 Best: 1, Worst: < 0.5




TP: True Positives, FP: False Positives, TN: True Negatives, FN: False Negatives, N: Number of samples



Test Set

Accuracy : 0.7351 [TP / N] Proportion of predicted labels that match the true labels. Best: 1, Worst: 0
Precision: 0.6989 [TP / (TP + FP)] Not to label a negative sample as positive. Best: 1, Worst: 0
Recall : 0.6537 [TP / (TP + FN)] Find all the positive samples. Best: 1, Worst: 0
ROC AUC : 0.7241 Best: 1, Worst: < 0.5

TP: True Positives, FP: False Positives, TN: True Negatives, FN: False Negatives, N: Number of samples



Training Set

* Naive Bayes *




Accuracy : 0.5635 [TP / N] Proportion of predicted labels that match the true labels. Best: 1, Worst: 0

Precision: 0.4893 [TP / (TP + FP)] Not to label a negative sample as positive. Best: 1, Worst: 0

Recall : 0.4054 [TP / (TP + FN)] Find all the positive samples. Best: 1, Worst: 0

ROC AUC : 0.5438 Best: 1, Worst: < 0.5




TP: True Positives, FP: False Positives, TN: True Negatives, FN: False Negatives, N: Number of samples



Test Set

Accuracy : 0.5734 [TP / N] Proportion of predicted labels that match the true labels. Best: 1, Worst: 0
 Precision: 0.4933 [TP / (TP + FP)] Not to label a negative sample as positive. Best: 1, Worst: 0
 Recall : 0.4147 [TP / (TP + FN)] Find all the positive samples. Best: 1, Worst: 0
 ROC AUC : 0.5520 Best: 1, Worst: < 0.5

TP: True Positives, FP: False Positives, TN: True Negatives, FN: False Negatives, N: Number of samples



Training Set

* Boosting *


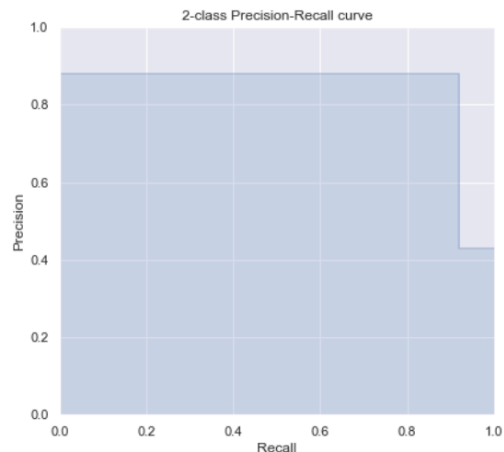

Accuracy : 0.9122 [TP / N] Proportion of predicted labels that match the true labels. Best: 1, Worst: 0

Precision: 0.8809 [TP / (TP + FP)] Not to label a negative sample as positive. Best: 1, Worst: 0

Recall : 0.9196 [TP / (TP + FN)] Find all the positive samples. Best: 1, Worst: 0

ROC AUC : 0.9131 Best: 1, Worst: < 0.5




TP: True Positives, FP: False Positives, TN: True Negatives, FN: False Negatives, N: Number of samples



Test Set

Accuracy : 0.9084 [TP / N] Proportion of predicted labels that match the true labels. Best: 1, Worst: 0
 Precision: 0.8662 [TP / (TP + FP)] Not to label a negative sample as positive. Best: 1, Worst: 0
 Recall : 0.9260 [TP / (TP + FN)] Find all the positive samples. Best: 1, Worst: 0
 ROC AUC : 0.9108 Best: 1, Worst: < 0.5

TP: True Positives, FP: False Positives, TN: True Negatives, FN: False Negatives, N: Number of samples



Summary

Training Models	Support Vector	Logistic Regression	Naïve Bayes	Boosting
Accuracy Scores (Training Sets)– 5 Text Columns	50.8%	49.3%	51.0%	50.1%
Accuracy Scores (Training Sets)– Identified 18 variables using forward feature selection	61.2%	74.6%	56.3%	91.2%

- ❖ Based on the data collected, we can use Machine Learning model — **Gradient Boosting** to help with the prediction of fake job listing. It provide the best accuracy at the rate of **91.2%**.
- ❖ AUC Score of **0.91** for both test and train test indicates a good classifier and overfitting is unlikely.
- ❖ Recruiters can use the predictive model to prevent such fraudulent job posters, improving the credibility of their recruitment domains.


- ❖ Try out other ML techniques and deep learning techniques to evaluate performance
- ❖ To deploy machine learning models to whole data frame. We could evaluate the accuracy, AUC score metrics if all the features were to be used as the predictor columns
- ❖ To present better models by using Grid Search to use the best estimators
- ❖ Resampling could have introduced bias. To test out variations of resampling with models
- ❖ To source for complementary and richer datasets from more recruitment sources to gather insightful data

THANKS!

Does anyone have any questions?

CREDITS: This presentation template was created by Slidesgo, including icons by Flaticon, and infographics & images by Freepik.

Please keep this slide for attribution.



RESOURCES

- ❖ <https://towardsdatascience.com/categorical-encoding-using-label-encoding-and-one-hot-encoder-911ef77fb5bd>
- ❖ <https://www.kdnuggets.com/2019/01/solve-90-nlp-problems-step-by-step-guide.html>
- ❖ <https://elitedatascience.com/imbalanced-classes>
- ❖ <https://www.cnbc.com/2020/10/06/job-scams-have-increased-during-the-covid-19-crisis-how-to-one.html>

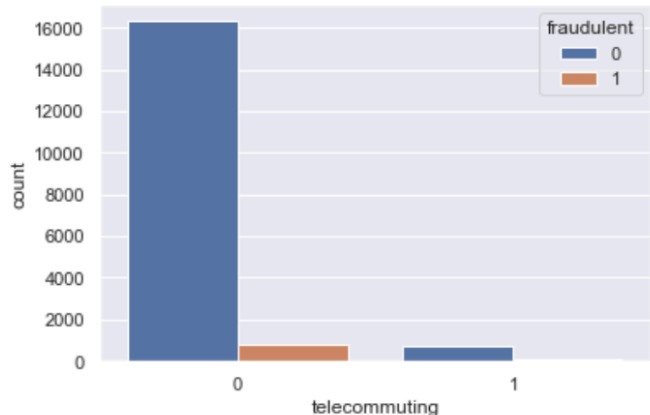



Appendices

Appendix

```
Job_postings['requirements']
```

```
0      Experience with content management systems a m...
1      What we expect from you:Your key responsibilit...
2      Implement pre-commissioning and commissioning ...
3      EDUCATION: Bachelor's or Master's in GIS, busi...
4      QUALIFICATIONS:RN license in the State of Texa...
...
17875   To ace this role you:Will eat comprehensive St...
17876   - B.A. or B.S. in Accounting- Desire to have f...
17877   At least 12 years professional experience.Abil...
17878   1. Must be fluent in the latest versions of Co...
17879   We want to hear from you if:You have an in-dep...
Name: requirements, Length: 17880, dtype: object
```



Encoding of categorical features

```
from sklearn.preprocessing import LabelEncoder
columns= ['telecommuting', 'has_company_logo', 'has_questions', 'employment_type',
          'required_experience', 'required_education', 'industry', 'function']
label_encoder = LabelEncoder()
for i in columns:
    Job_postings[i] = label_encoder .fit_transform(Job_postings[i])
```

Basic text cleaning to each text column

```
def clean_text(company_profile):
    # reduce multiple spaces and newlines to only one
    company_profile = re.sub(r'(\s+|\n\n+)', r'\1', company_profile)
    # remove double quotes
    company_profile = re.sub(r'\"', '', company_profile)
    # remove extra whitespace and special characters
    company_profile = re.sub(r'\s+$', '', company_profile)
    company_profile = re.sub(r'^^+$', '', company_profile)
    company_profile = re.sub(r'^a-zA-Z\s+', '', company_profile)
    #lowercase
    company_profile= company_profile.lower()
    # remove text between square brackets
    company_profile =re.sub('\[[^\]]*\]', '', company_profile)

    # removes punctuation
    company_profile= re.sub(r'[\W\s]', '', company_profile)
    company_profile = re.sub(r'\n', '', company_profile)
    company_profile= re.sub(r'[0-9]+', '', company_profile)
    company_profile= re.sub(r'[0-9]+', '', company_profile)
    # remove URLs
    company_profile = re.sub(r'https://t.co\S+\s*', '', company_profile)
    company_profile = re.sub(r'http://t.co\S+\s*', '', company_profile)

    return company_profile
```

Feature Extraction

```
Job_postings['word_count_r'] = Job_postings['requirements'].apply(lambda y: len(str(y).split(" ")))
Job_postings[['requirements', 'word_count_r']].head()
```

	requirements	word_count_r
5	experience with content management systems a m...	115
6	what we expect from youyour key responsibility...	187
7	implement precommissioning and commissioning p...	164
8	education bachelors or masters in gis business...	174
9	qualificationsrn license in the state of texas...	89

Further text cleaning

```
def convert_text(description, remove_stop=True, lemma_words=False):  
    # Remove stop words  
    if remove_stop:  
        description = description.split()  
        description = [w for w in description if not w in stop]  
        description = " ".join(description)  
  
    #remove lemma_words  
    if lemma_words:  
        description = description.split()  
        wl = nltk.stem.WordNetLemmatizer()  
        lemma = ' '.join([wl.lemmatize(word) for word in description.split()])  
    # tokenize  
    description = nltk.word_tokenize(description)  
  
    return description
```

```
Job_postings.head()
```

ice	...	char_count_r	char_count_b	char_count_t	char_count_cp	word_density_d	word_density_cp	word_density_r	word_density_t	word_density_b	description1
5	...	715	0	15	710	6.064516	5.035461	6.217391	7.500000	0.000000	foodfast growingj ames...
7	...	1172	1022	35	909	5.333333	6.060000	6.267380	5.833333	4.542222	organise dfocuse dvibra...
0	...	1173	0	34	728	5.980000	5.352941	7.152439	8.500000	0.000000	clientlo catedho uston...
6	...	1228	669	28	509	6.414493	5.988235	7.057471	5.600000	6.968750	company esrienvi ronmen...
6	...	653	19	17	1327	7.142857	6.473171	7.337079	5.666667	6.333333	jobtitlei temizati onr...

Count Vectorization

```
X = Job_postings['title1']
y= fraudulent_upsampled['fraudulent'].sample(n=17880)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2, random_state = 42)
# Learn a vocabulary dictionary of all tokens in the raw documents
counts.fit(X)

# Transform documents to document-term matrix.
X_train_count_1 = counts.transform(X_train)
X_test_count_1 = counts.transform(X_test)

print(X_train_count_1.shape)
print(X_test_count_1.shape)
print(y_train.shape)
print(y_test.shape)
```


```
(14304, 26)
(3576, 26)
(14304,)
(3576,)
```

```
df_train_t=pd.DataFrame(X_train_count_1.todense(),columns=counts.get_feature_names())
df_test_t=pd.DataFrame(X_test_count_1.todense(),columns=counts.get_feature_names())
```


```
text_train_1 =pd.concat([df_train_d, df_train_t],axis=1)
text_test_1 =pd.concat([df_test_d, df_test_t],axis=1)
```

Confusion Matrices


SVM



Logistic Regression



Naïve Bayes



Boosting

